

Web Protocols

Resources

Request and Responses

Error Codes



Hypertext Transfer Protocol

- HTTP is an application/layer protocol on server port 80
- was created in 1993 by Tim Berners-Lee, a physicist

"The Hypertext Transfer Protocol (HTTP) is an application-level protocol with the lightness and speed necessary for distributed, collaborative, hypermedia information systems." RFC1945
- is one of the simplest protocols
 - Writing handler code is really easy
- has only two phases:
 - Request (Client → Server)
 - Response (Server → Client)
- is Stateless (as originally defined)
- Some REQ/RESP are idempotent, so optimal to cache
 - Every transaction has no memory of previous transactions

Hypertext Transfer Protocol

- HTTP had 4 main versions
 - v0.9. Original release and simplest
 - v1.0 HTTP Sessions introduced as a series of request/response pairs
 - v1.1 Added partial downloads, better caching
 - v2 Significant upgrade, introduced in 2014
- Numerous optimisations have been added, but the basic protocol still works
 - Main goal in these optimisations was to improve Performance and reduce Latency
 - Generally only ver. 1.1 and 2 are now used.

Resource Retrieval using HTTP

Retrieval with the HTTP protocol has five main steps, usually implemented by the browser as client:

1. Map the server domain name to an IP address
2. Establish a TCP/IP connection to the web server
3. Transmit a **request** to the server, which includes a request method, and any additional information
4. Receive a **response** from the server, such as HTML text, an image, or other information
5. Server terminates the TCP/IP Connection (if timeout, then client terminates)

HTTP Example

Using elnet

Normally telnet	↔	telnetd	port 23
ssh	↔	sshd	port 22
http	↔	httpd	port 80

```
% telnet titan.csit.rmit.edu.au 80
```

```
Trying 131.170.5.131...
```

```
Connected to titan.csit.rmit.edu.au.
```

```
Escape character is '^['.
```

```
GET /
```

← Type a blank Line (Two CRs in a row)
----- No HTTP Request Body -----

```
HTTP/1.1 200 OK
```

```
Date: Fri, 7 Feb 2018 17:35:20 GMT
```

```
Server: Apache/2.4.6 (Red Hat Enterprise Linux)
```

```
Last-Modified: Thu, 14 Jan 2016 02:26:01 GMT
```

```
ETag: "76-5294201ff0a4e"
```

```
Accept-Ranges: bytes
```

```
Content-Length: 118
```

```
Vary: Accept-Encoding
```

```
Connection: close
```

```
Content-Type: text/html
```

HTTP Response Header

← blank line between HTTP
Response Header and content

```
<html>
```

```
<body>
```

```
  <iframe name="my_iframe" src="motd.txt" width=800 height=600  
    align="centre"></frame>
```

```
</body>
```

```
</html>
```

```
Connection closed by foreign host.
```

Statelessness of HTTP

HTTP is a **stateless** protocol,

- there no memory involved in the transaction – all are alike
- so each HTTP transaction between a client and a server is an independent entity. That is, in our example, the request for the document body and the image are *unrelated* transactions.
- Does not use Session-layer state management
- With large numbers of transactions arriving at a web server, statelessness brings efficiency benefits, since a server does not need to track clients, sessions, or histories.

However if state is required, for example if a site uses a shopping cart, then other methods have to be implemented, for example cookies.

Idempotence

- An operation in a server is *idempotent*, if that operation does not change the state of the server.
- This means:
 - You can do the operation many times and the result is always the same
 - You can change how you do it without affecting the server
 - You can test it without changing what clients see
 - You can assume what the result will be if you did it before (recently).
 - You can cache the operation
 - you can pretend to do it but really just return the result as if you had done it.
 - It allows all forms of virtualisation to be simplified
- Idempotence is an important performance NFA

HTTP Request & Response messages

- Request Header
 - Request line, starting with a Request Method
 - eg GET /path/to/resource
 - Other header lines
 - Blank line to separate header from content
 - Content
 - data to send to server (usually empty for GET)
- Response
 - Response line
 - return protocol + error code (eg HTTP 1.1 200 Ok)
 - Other header lines
 - Blank line to separate header from content
 - Content
 - returned data from server

HTTP Request Methods

- On the request line, there are 10 possible request methods that can be made to a HTTP server:

Command	Action
GET*	Return Document Contents
HEAD*	Return the document header only (Metadata)
PUT*	Replace the document with the following data
POST	Treat the document as a script and send data
OPTIONS*	Return a list of methods that work on this URI
CONNECT	Set up a TCP connection through HTTP (eg TLS)
Others:	CREATE, DELETE, TRACE*, PATCH

* Idempotent methods

HTTP Get Method

URL: `http://server.com/path/to/resource?opt1=A&opt2=B`

- GET retrieves a resource from the web server.
- GET can include additional information that will help the web server formulate a response.
- Format:
GET /path/to/resource?options HTTP-version
- Examples
 - GET /index.html
 - Get file index.html from the root directory
 - GET /index.html HTTP/1.1
 - As above using HTTP/1.1 version
 - GET /maps/showmap?lat=39&long=144&zoom=3
 - Get a map, using ?options to determine location and size

HTTP Head Method

URL: `http://server.com/path/to/resource?opt1=A&opt2=B`

- HEAD does NOT retrieve a resource from the web server.
- Instead it sets the same metadata that a GET would have set, but does not retrieve the resource
- Used for checking whether a server copy has been updated since last download.
- Format:
`HEAD /path/to/resource?options HTTP-version`
- Examples
`HEAD /index.html HTTP/1.1`
 - Like GET but without download resource

HTTP Put Method

URL: `http://server.com/path/to/resource`

- PUT sends a resource to the web server.
- typically used for file upload
- has security issues (e.g. XSS)
- Format:
PUT /path/to/put/resource HTTP-version
- Examples
PUT /index.html HTTP/1.1
 - As above using HTTP/1.1 versionPUT /logs/locationdata.dat?lat=39&long=144&id=54642
 - Put log data, using ?options to determine location and ID

HTTP Post Method

URL: `http://server.com/path/to/resource/resource?opt1=A&opt2=B`

- POST sends data to the web server.
- typically used for data entry forms
- similar to GET except request content is not blank
 - instead it contains the options, that in GET would have been appended to the URL
 - can be compressed and/or encrypted, unlike HTTP headers.
- Format:
**POST /path/to/put/resource HTTP-version
opt1=A
opt2=B**
- **Example**
POST /logs/locationdata.dat HTTP/1.1
lat=39
long=144
id=54642
<blanks>
 - Put log data, using ?options to determine location and ID

Selected HTTP Status Codes

Code	Text	Description
<u>2xx</u>		<u>Success</u>
200	OK	URL found, Data Follows
202	Accepted	Request accepted for later processing
204	No Response	OK, but no data follows
206	Partial Content	GET partially succeeded (eg a PDF page)
<u>3xx</u>		<u>Redirection</u>
301	Moved	The URL has moved permanently
302	Found	The URL has moved temporarily
<u>4xx</u>		<u>Client Errors</u>
400	Bad Request	Syntax Error
401	Unauthorised	Authorisation Error
403	Forbidden	URL can never be retrieved
404	Not found	URL is not here
<u>5xx</u>		<u>Server Errors</u>
500	Internal Error	Something weird happened
501	Not Implemented	Feature not supported

HTTP/1.1

SPDY

Google was pushing SPDY, but now it is deprecated since May 2016

- SPDY ('speedy') acts as a tunnel/filter for HTTP. It:
 - Improves security
 - Usually requires SSL/TLS, but can work without it
 - Reduces page load latency
 - Compresses data stream & headers
 - (GZIP or DEFLATE)
 - Prioritises the data items sent (*Prioritised parallelism*)
 - items not necessarily sent in document order
 - Minimises the number of separate connections
 - Implements Server Push
 - Create an *imagined* (*virtual*) request, then respond to it.

**SPDY sits in the
Application Layer
so it cannot be
invisible and
browsers need to
know about it.**

HTTP/2

- HTTP/2 (*Not HTTP 2.0, or HTTP/2.0*)
 - Complete superset of HTTP/1.1 (so fully supports it)
 - High level syntax is the same as HTTP/1.1
 - Low level signalling and transports are different
 - Based on SPDY, so same benefits as it
 - **Content Negotiation** mechanism (HTTP/1.1 /2 or other)
 - ALPN = Application Layer Protocol Negotiation
 - Improve page load latency like SPDY (with minor differences)
 - Header Compression, Server Push, Pipelining, Multiplexing
 - Fix "head of line" blocking issue
 - Encryption (defined for HTTP/HTTPS)
- The HTTP/2 standard is available from:
 - <http://www.ietf.org/rfc/rfc7540.txt> (2015)

SPDY sits in the Application Layer so it cannot be invisible and browsers need to know about it.

Criticisms of HTTP/2

- Flow Control
 - Handled by TCP as SYN/ACK pairs on separate connections
 - HTTP/2 Duplicates this using stream ID on a single connection
 - Breaks the protocol layering principles of TCP/IP
- Encryption
 - Some in Standards Committee tried to force encryption. Failed.
 - But browsers only use HTTP/2 over HTTPS, so de-facto succeeded
- HTTP/2 flow control is explained in detail at
 - <http://qimate.com/what-is-multiplexing-in-http2/#prettyPhoto>

QUIC

(Quick UDP Internet Connections) - Google


- aims to replace TCP
- reduce handshaking
 - TCP: 3 at the front, 2 at the end
- reduce RTT (round-trip-time)
 - by reducing number of RTs
- SSL/TLS-like security over UDP
- Allows stream multiplexing
 - (multi requests, receive responses in any order)
 - like SPDY
- Application layer congestion control which handles multiplexed connections
 - unlike TCP, which does not
- Fall back to TCP
 - if UDP is blocked

QUIC sits in the Transport Layer so it does not affect HTTP/1.1, but modifies signals for speed.

New Framework

- Overall design is to **improve network latency**
 - improve by reducing **number of packets** sent
 - Header and body compression / encryption
 - Reduce duplication of headers and other information
 - Multiplex many parallel streams and manage them

more: <https://ma.ttias.be/googles-quic-protocol-moving-web-tcp-udp/>



Observations of New Framework

- These last few years showed how a major standard protocol can evolve.
 - Application Layer (a new HTTP?)
 - HTTP/2 is mostly HTTP/1.1 and SPDY combined
 - SPDY was temporary and is now implemented in HTTP/2
 - Transport Layer (a new TCP?)
 - QUIC is UDP with TLS, flow control and compression
 - QUIC is temporary and if no major issues arise, may well be incorporated into the next version of TCP

Summary

- HTTP
 - Definition
 - Resource Retrieval
 - Statelessness and Idempotence
 - Request and Response model
 - Request Methods and Error Codes
 - HTTP/2 / SPDY / QUIC
 - New Web Framework